

## CS221 Project Progress Report

Our project task is to predict binding affinities of modified bis-guanidinium neurotoxins (**the ligand**) to different isoforms of voltage-gated sodium ion channels (**the protein**). To review, our training data is a set of 300 experimentally collected binding affinities of ligands to proteins. An example datapoint is as follows:

Toxin Subclass	Compound (Common Name)	Compound (SMILES)	Protein	Protein DEKA Sequence	Binding Affinity (IC50)	Unit	pKi	Ki	Potent	Source
STX	STX.N21.1	<chem>[NH2+]=C1N[C@@H](CO C(NC)=O)[C@H](NC2=[NH2+])[C@]3(N2)N1CCC3(O)O</chem>	rNaV 1.4	TQDYW EN CGEWI ET TFKGW MD TSAGW DG	1.9	nM	0.02228	0.95	1	RTT

“Toxin subclass” refers to the scaffold on which the ligand was created. There are four scaffolds in our dataset (STX, ZTX, GTX, TTX); each is a different bis-guanidinium natural product. The “Compound (Common Name)” column stores the ID of the ligand. It maps to a separate file with 2-D line drawings of the ligand for easy visualization. There are about 40 unique ligands in our dataset. The “Compound (SMILES)” column stores a string representation of the ligand. The SMILES string encodes all the information necessary to visualize the ligand in three dimensions (such as the atomic composition, bonding and structure). The “Protein” column stores the common name of the protein. All proteins in this dataset are isoforms or mutated isoforms of rat or human sodium channels; there are about 70 unique proteins in our dataset. Because all our ligands are based on bis-guanidinium neurotoxins that are hypothesized to restrict the access of sodium ions to the protein by plugging up the protein above the sodium selectivity filter like a cork, modifications made to the protein sequence are concentrated in the DEKA region, a loop of amino acids bordering the point in the protein channel where toxins are hypothesized to dock. “Protein DEKA Sequence” stores the amino acids present in this region; the composition of amino acids in this region uniquely identifies each of the 70 proteins in our dataset. “Binding affinity” stores the experimentally measured IC50 in nanomolars of the specific ligand/protein pair identified by the row. “pKi” and “Ki” store the inhibition constant, another measure of binding affinity. Although this value can be experimentally measured, it was not in our case and was rather calculated by making the rough approximation that  $K_i = \frac{1}{2} \cdot (IC_{50})$  for the case of competitive protein-ligand interactions. Since all ligands in our dataset competitively restrict the access of sodium ions into the protein, this approximation is a fair assumption to make. “Potent” is a binary 0/1 label; ligands are deemed potent if their IC50 is less than or equal to 500 nM. Finally, “Source” maps to the researcher that collected the experimental datapoint; all datapoints were gathered by hand by this project group from theses written by members of the Du Bois group at Stanford.

To summarize, our problem input is a protein/ligand pair, and our desired output is a measure of binding affinity. The algorithm/model we would like to use to solve our problem is linear regression passed along layers of a pre-trained convolutional neural network called Pafnucy. Pafnucy [paphnusus] is a 3D convolutional neural network developed by Marta Stepniewska-Dziubinska that predicts binding affinity for protein-ligand complexes. It was trained on the PDBbind database and tested on the CASF “scoring power” benchmark.

Protein-ligand complexes must be transformed and encoded into tensors in order to be utilized by a neural network. Pafnucy accomplishes this by cropping three-dimensional coordinate files of a protein-ligand complex to a defined 20-Å cubic box centered at the geometric center of the ligand. Positions of heavy atoms are discretized using a 3D grid with 1-Å resolution, allowing for each input to be represented as a 4D tensor in which each point (atom) is defined by its Cartesian coordinates (the first three dimensions of the tensor) and a vector of features (the last dimension). Inputs to Pafnucy use the following 19 features in its last dimension to describe an atom:

- 9 bits (one-hot or all null) encoding atom type: B, C, N, O, P, S, Se, halogen and metal
- 1 integer (1, 2, or 3) storing atom hybridization.
- 1 integer storing the number of bonds with heavy atoms.
- 1 integer storing the number of bonds with heteroatoms.
- 5 bits encoding properties like hydrophobicity, aromaticity, H-bond acceptor/donor status, and ring presence.
- 1 float storing partial charge.
- 1 integer (1 for ligand, -1 for protein) to distinguish between the two molecules.

Pafnucy was trained and tested on protein-ligand complexes from the PDBbind database v. 2016. This database contained 3D structures of molecular complexes and their corresponding binding affinities expressed with pKa (-logKi) values. All complexes were protonated and charged using UCSF Chimera with Amber ff14SB for standard residues and AM1-BCC for non-standard residues and ligands. About 11,906 complexes total were used in Pafnucy's training set; 1000 complexes were used in validation, and 290 complexes were used as an external test set. The training, test, and validation sets were all disjoint.


Pafnucy's architecture is a deep convolutional neural network with a single output neuron for predicting binding affinity. The input is first processed by a block of 3D convolutional layers. Pafnucy uses 3 convolutional layers with 64, 128, and 256 filters. Each layer's filters have 5-Å cubic resolution and are followed by a max pooling layer with a 2-Å cubic patch. The result of the last convolutional layer is flattened and used as input for a block of dense layers. Pafnucy uses 3 dense (fully connected) layers with 1000, 500, and 200 neurons. Dropout with drop probability of 0.5 was used for all dense layers. Both convolutional and dense layers are composed of ReLU units.

In the training stage, Pafnucy was initialized with values for the convolutional filter weights drawn from a truncated normal distribution with mean 0 and standard deviation 0.001. Biases were set to 0.1. The weights in the dense were initialized with a truncated normal distribution with mean 0 and standard deviation  $1/\sqrt{n}$ , where n was the number of incoming neurons for a given layer. The biases for the dense layers were set to 1.0. The network was trained using the Adam optimizer with a learning rate of  $10^{-5}$  and 5 examples per minibatch. Training was carried out for 20 epochs, but the model performance was optimized after 14 epochs of training. Dropout and L2 weight decay ( $\lambda = 0.001$ ) was implemented to reduce overfitting. Every input structure was presented to the network in 24 different orientations to develop a model that is not sensitive to ligand-receptor complex orientation. Thus, 24 training examples are generated for each protein-ligand complex.

We want to perform transfer learning on Pafnucy using the dataset of 300 neurotoxin-sodium channel binding pairs. Our process of featurizing the dataset we had gathered to create input for Pafnucy is as follows:

First, we wrote a script to convert the ligands SMILES strings to 3D structure files with the .pdb (Protein Data Bank) file extension using the online SMILES translator developed by the NCI (<https://cactus.nci.nih.gov/translate/>):

[Home](#) | [About](#) | [Contact](#) | [Disclaimer](#) | [Privacy](#) | [Notes](#)


**NATIONAL CANCER INSTITUTE**

NCI/CADD Group

### Online SMILES Translator and Structure File Generator

[Form](#) | [News](#) | [Help](#) | [Acknowledgments](#)

Input Format	Unique SMILES Output Format (Output SMILES)
<input type="text" value="C12C3C4C1C5C4C3C25"/> Start Structure Editor Please choose this field if you want to submit your own SMILES strings or create a SMILES string using the Structure Editor. A submitted file has precedence, so delete any entry below if you want to submit a new SMILES string.	<input checked="" type="radio"/> Display on screen <input type="radio"/> SMILES TXT file <input type="radio"/> SDF <input type="radio"/> PDB <input type="radio"/> MOL (only single structure generated)
<input type="button" value="Choose File"/> No file chosen Please choose this field if you want to translate your own files. The service will automatically recognize SD files (single and multiple structure), text files with multiple SMILES fields, MOL files and PDB files (and in fact any other format CACTVS recognizes).	Use <input type="radio"/> Kekule or <input checked="" type="radio"/> Aromatic SMILES representation (choose "Aromatic" for closer approximation to Daylight USMILES)  SD, PDB or MOL files should contain <input type="radio"/> 2D <input checked="" type="radio"/> 3D coordinates
If the input file contains a single structure, the output will also be single structure. Multiple structure input formats will generate multiple structure output for those formats that support this. Otherwise, only the first structure will be used. SD files will contain a UNIQUE SMILES field for unique SMILES and an USER_SUPPLIED_SMILES field for the user-supplied SMILES (if available)	

Next, we generated .pdb files for each one of our 70 unique sodium channels. Starting from a cryo-EM image of human sodium channel 1.7 (hNav1.7) bound to saxitoxin with 3.2-Å resolution, we extracted the coordinates of the protein and

identified the amino acids forming the DEKA loop. Next, we mapped hNaV1.7's DEKA loop onto rat sodium channel 1.4's DEKA loop by feeding their respective amino acid sequences to BLAST (the Basic Local Alignment Search Tool). We also mapped each of the rat sodium channels 1.1-1.9 to the corresponding human isoform.

Having acquired the sequence correspondence code, we wrote a script to edit the .pdb file of hNaV1.7. Our script modifies amino acids in the DEKA loop by editing the coordinates and atomic identities of the relevant sidechains. The figure below shows PyMol's visualization of the structure file of hNaV1.7 (our template). The DEKA region is in the center of the protein. To the right is a text representation of the same structure file with the amino acid at position 1410 highlighted. The regions in green correspond to the backbone atoms, which remain the same for each amino acid. The regions in red represent the side chain atoms, which change depending on amino acid identity. Side chains are attached at the atom labeled 'CA' in the third column of the green region; thus, by treating the CA atom as the zero-coordinate and adding/removing atoms from the red region as appropriate, we can edit the PDB file to generate three-dimensional structure files of the proteins for which we don't have cryo-EM images. We generated theoretical three-dimensional coordinates for new side chains by referencing the geometry of amino acids in vacuo.



ATOM	6308	CH2	TRP	A1408	142.149	137.344	137.799	1.00	23.41	C
ATOM	6309	N	THR	A1409	143.099	133.510	142.648	1.00	25.11	N
ATOM	6310	CA	THR	A1409	143.746	134.492	143.509	1.00	25.11	C
ATOM	6311	C	THR	A1409	144.597	133.916	144.616	1.00	25.11	C
ATOM	6312	O	THR	A1409	145.611	134.489	144.949	1.00	25.11	O
ATOM	6313	CB	THR	A1409	142.792	135.567	144.054	1.00	25.11	C
ATOM	6314	CG2	THR	A1409	141.668	135.824	143.087	1.00	25.11	C
ATOM	6315	OG1	THR	A1409	142.261	135.149	145.310	1.00	25.11	O
ATOM	6316	N	ILE	A1410	144.228	132.770	145.231	1.00	23.64	N
ATOM	6317	CA	ILE	A1410	145.044	132.197	146.298	1.00	23.64	C
ATOM	6318	C	ILE	A1410	146.408	131.844	145.738	1.00	23.64	C
ATOM	6319	O	ILE	A1410	147.412	132.039	146.396	1.00	23.64	O
ATOM	6320	CB	ILE	A1410	144.434	130.932	146.914	1.00	23.64	C
ATOM	6321	CG1	ILE	A1410	142.951	131.125	147.188	1.00	23.64	C
ATOM	6322	CG2	ILE	A1410	145.154	130.580	148.202	1.00	23.64	C
ATOM	6323	CD1	ILE	A1410	142.208	129.819	147.316	1.00	23.64	C
ATOM	6324	N	ILE	A1411	146.405	131.289	144.526	1.00	23.81	N
ATOM	6325	CA	ILE	A1411	147.650	130.925	143.861	1.00	23.81	C
ATOM	6326	C	ILE	A1411	148.541	132.125	143.545	1.00	23.81	C
ATOM	6327	O	ILE	A1411	149.731	132.078	143.771	1.00	23.81	O

Having created three-dimensional structure files for ligands and proteins, we proceeded to dock the ligands against the proteins to create ligand-receptor complex inputs for Pafnucy. To do so, we used Vina, the current industry standard for protein-ligand docking. Vina uses molecular mechanics simulations combined with multithreading to predict the geometry and binding affinity of small molecule ligands to protein pockets. To use Vina, we first wrote scripts to convert the PDB files of our ligands and proteins to PDBQT files using Vina's companion software, AutoDock. The PDB->PDBQT conversion process adds polar hydrogen atoms and partial charges to the protein and ligand structure files. Having generated PDBQT files for all our protein and ligand pairs, we wrote a script to push the data through Vina. Vina takes a configuration file that specifies the search space for the docking program, which we defined as a 20-Å cubic box oriented at the geometric center of the DEKA loop. Vina's output is a pdb file of ligand docking poses along with a predicted binding affinity (in kcal/mol). We selected the highest affinity result as our hypothesized docking pose and stored the binding affinity that Vina output as our oracle.

Our final step was to overlay the protein structure files over our ligand docking file, to combine the two objects, and to extract a 20-Å cubic box from the protein, centered at the geometric center of the ligand, to offer to Pafnucy as input. We accomplished this by writing a Python script in PyMol. We ran the pocket structure file and ligand structure file through Pafnucy's featurizer to create the 4D tensors described previously. Then, we fed the input to Pafnucy to output numbers for predicted binding affinity.

Our project baseline was a unigram naive bayes classifier using SMILES strings and protein amino acid sequences from the DEKA loop as features, which achieved a classification accuracy of 47% on our training dataset (using an 80:20 train/test split). Note that we trained a classification, not a regression algorithm for our baseline, so the baseline accuracy cannot be directly compared against the oracle accuracy as is. Our project oracle is Vina's binding affinity output, which we converted to  $K_i$  using the following equation, where  $\Delta G$  is Vina's output,  $T=298K$ , and  $R=1.986$  cal/mol K.

$$K_i = \exp(\Delta G/RT)$$

We have yet to fully automate this process for our entire dataset, but we were able to run our baseline algorithm for a small subset of the data.

The oracle (Vina) achieved a percent error (defined as  $100\% \times [\text{observed binding affinity} - \text{expected binding affinity}] / \text{expected binding affinity}$ ) of -302%.

In comparison, our initial algorithm (a single run through Pafnucy) achieved a percent error of -327%.

The output generated by Pafnucy compares favorably to the output generated by the oracle (they match almost exactly). This is good because it suggests that our initial algorithm is already very well-adapted towards solving this problem, and it is a good template from which to optimize further. However, both Pafnucy and the oracle's generated binding affinity predictions are far off from the actual experimental  $K_i$ 's. This may actually be a problem with our data translation process. Recall that we were able to glean only  $IC_{50}$ 's from our initial dataset, which we translated to  $K_i$ 's via a rough approximation that  $K_i = IC_{50}/2$ . The actual conversion process involves a more rigorous consideration of Michaelis-Menten kinetics, a calculation that depends on the concentration of the protein and ligands in solution. We will work on gathering the data to perform this calculation as we continue to develop this project; hopefully, this will correct our values moving forward.